



ALEX DIMMOCK

DATA ENGINEER

atedimmock@gmail.com (mailto:atedimmock@gmail.com)

Milton Keynes, England, United Kingdom

 [Alex - LinkedIn](https://www.linkedin.com/in/alex-dimmock/) (https://www.linkedin.com/in/alex-dimmock/)

 [Alex - GitHub](https://github.com/AtomicAlexD) (https://github.com/AtomicAlexD)

Data Engineer & Architect: Building Next-Gen Data Platforms and AI-Powered Solutions

I architect scalable data platforms and production AI systems, solving complex data challenges with Databricks, Azure, and the modern data stack.

Proven expertise in agentic workflows and RAG applications. I've shipped production systems using Pydantic AI, DSPy, and LangChain/LangGraph with MLflow for model and agent tracking at scale.

Experienced with Kimball dimensional modelling, data vault, and enterprise architectures deployed across global organizations.

Strong track record designing and implementing streaming pipelines across multiple technologies and platforms.

Passionate about building and mentoring high-performing data engineering teams, I've had the privilege of helping talented engineers launch successful careers.

WORK EXPERIENCE

Senior Data Engineer at [The Access Group](#)

(<https://www.theaccessgroup.com/>)

November 2022 - Present

At the forefront of architecting and developing a greenfield Data platform to cater to the entire business on a global level. Leveraging a tech stack that includes Databricks, Azure, Apache Spark, Python, and SQL to build next-generation AI-powered data solutions.

- Spearheading the design and architecture of the Data platform on Databricks, ensuring scalability, reliability, and performance
- Involved in every step of planning from data governance, data architecture down to build and initialization of custom Python framework for data engineering on the platform
- Leveraging Spark to full capabilities, utilizing DLT pipelines in Databricks and exploring all platform functionality
- Designed main ingestion pipelines that take federated catalogs and load them into platform with minimal latency
- Built custom Python application that enables other data engineers across the business to work more efficiently
- Collaborating cross-functionally with group level data teams and product managers to understand business requirements and translate them into scalable data solutions
- Fostering a culture of innovation and data-driven decision-making across the organization

Head of Data at adam (HTT)

September 2020 - November 2022

Led the transformation of the company's data capabilities during a period of rapid growth. Formed and managed a dedicated team of Data Engineers and Data Analysts to revolutionize how data was harnessed and presented, both internally and externally.

- Promoted to Head of Data to align data capabilities with modern data stack during company's rapid growth
- Built and led dedicated team of Data Engineers and Data Analysts
- Led complete reconstruction of internal data warehouse grounded in Kimball methodology
- Planned and implemented each phase including comprehensive data modelling
- Designed architecture leveraging Azure technologies - Power BI front end, Azure SQL data warehouse, and Azure Data Factory for ELT
- Re-evaluated and remodelled all existing data sources from internal systems and main application
- Completed four pivotal phases of delivery, propelling data capabilities to forefront of organizational operations
- Enabled more efficient work processes, leveraged data-driven insights, and improved decision-making across the company

Senior Database Developer at adam (HTT)

February 2018 - September 2020

Promoted to senior role with the opportunity to assemble and nurture a high-performing team. Focused on delivering procurement services to local government entities while mentoring team members and addressing comprehensive business intelligence and reporting needs.

- Promoted to senior role following company division
- Assembled and nurtured high-performing team, mentoring individuals from internal support roles into accomplished data professionals
- Worked closely with key stakeholders across the organization
- Addressed comprehensive business intelligence and reporting needs
- Supported finance department with monthly reporting
- Maintained legacy data warehouse
- Successfully managed second infrastructure migration from one private cloud to another
- Earned Microsoft MCSE for Data Management and Analytics during tenure

Database Developer at Dynamic Procurement Holdings Ltd

May 2016 - February 2018

Served as the inaugural and sole data resource at DPH, assuming full responsibility for the company's data estate including legacy data warehouse and on-site infrastructure. Led critical infrastructure migration to private cloud environment.

- First and sole data resource at DPH with full responsibility for company's data estate
- Managed legacy data warehouse and on-site infrastructure
- Led critical infrastructure migration project transitioning on-site infrastructure to private cloud
- Performed various database administration activities
- Implemented availability groups and upgraded to SQL 2014
- Delivered significant enhancements to existing systems

Junior BI Consultant at Dynamic-BI

May 2012 - May 2016

Started as an apprentice and progressed to business intelligence consultant. Gained comprehensive experience in the Microsoft data stack, managed on-site infrastructure, and delivered end-to-end data warehousing solutions to clients nationally and internationally.

- Earned CompTIA certifications as an apprentice
- Managed all on-site infrastructure including phone systems, network devices, and servers

- Performed database administration for Microsoft SQL and Oracle servers vital for development and QA
- Shadowed senior business intelligence consultants to gain proficiency in Microsoft data stack (SSIS, SSRS, SSAS, T-SQL)
- Actively participated in stakeholder meetings, developing soft skills alongside technical expertise
- Progressed to business intelligence consultancy role
- Collaborated with senior colleagues to strategize, build, and deliver end-to-end data warehousing solutions
- Worked with clients nationally and internationally
- Traveled to USA to deliver comprehensive financial BI solution to one of the world's largest grain suppliers

PROJECTS

Next-Generation AI-Powered Data Platform

Leading development of The Access Group's next-gen data platform on Databricks, enabling AI-powered insights and supporting global business strategy at scale.

- Architected and built greenfield Databricks platform serving global business operations with production-grade reliability and minimal latency ingestion
- Designed and implemented federated data ingestion pipelines leveraging Spark and DLT, consolidating multi-source data at scale
- Pioneered production agentic AI workflows using Pydantic AI, DSPy, and LangChain/LangGraph with MLflow for experiment tracking, model registry, and deployment governance
- Built standardized MLflow pipelines enabling reproducible model training, versioning, and A/B testing across the organization
- Built custom Python framework for data engineering that standardized development practices and accelerated delivery for 10+ engineers
- Established data governance and architecture standards enabling self-service analytics for global teams

Enterprise Data Warehouse Modernization

Transformed adam (HTT)'s data capabilities by rebuilding their entire data warehouse—replaced legacy systems with modern cloud architecture, enabling analytics teams to operate independently.

- Designed and implemented dimensional data warehouse using Kimball methodology, modernizing legacy infrastructure
- Built and mentored dedicated team of Data Engineers and Data Analysts from ground up

- Migrated and remodeled 100+ data sources from legacy systems into new architecture
- Deployed Azure Data Factory → Azure SQL data warehouse → Power BI pipeline, enabling real-time analytics and self-service reporting
- Delivered 4 major phases on schedule, enabling data-driven decision-making across the organization

Cloud Infrastructure Migrations

Orchestrated large-scale infrastructure migrations at Dynamic Procurement Holdings, successfully transitioning mission-critical systems to cloud while maintaining operational continuity.

- Led two critical infrastructure migrations—on-premise to private cloud, then between private cloud providers
- Implemented SQL Server Availability Groups and upgraded legacy systems (SQL 2012 → 2014) with zero downtime
- Designed migration strategies that maintained business continuity while modernizing the data estate

International Financial Analytics Platform

Delivered high-impact BI solution to global Fortune 500 organization, establishing complex financial analytics capabilities for international operations.

- Designed and delivered end-to-end financial data warehouse for one of the world's largest grain suppliers
- Built comprehensive BI solution supporting multi-currency financial reporting and executive dashboards
- Managed international engagement including on-site delivery in USA

SKILLS

SQL, Python, PySpark

Expert

Databricks, Apache Spark, Data Streaming, Data Architecture

Expert

Azure, Data Modeling, Dimensional Modeling, Data Vault

Advanced

MLflow, Model Registry, Experiment Tracking, Machine Learning

Advanced

Agentic Workflows, RAG Applications, LangChain, Pydantic AI, DSPy

Advanced

Azure Data Factory, Power BI, Microsoft Fabric

Proficient

INTERESTS

AI and Machine Learning

As a Data/AI Engineer, I'm passionate about building RAG-based applications and agentic workflows. I enjoy exploring new AI toolsets and finding innovative ways to integrate artificial intelligence into data platforms to solve complex business problems.

Staying current with the latest developments in AI and machine learning allows me to design solutions that are not just data-driven, but intelligently automated and forward-thinking.

Team Building and Mentorship

One of my greatest professional satisfactions comes from building and nurturing teams in the data engineering space. I take pride in helping talented individuals develop their careers and grow into accomplished data professionals.

Throughout my career, I've had the privilege of mentoring people from various backgrounds, transforming them into skilled data engineers and analysts. Seeing their success and career progression is incredibly rewarding.

Data Architecture and Modern Data Stack

I'm deeply interested in data architecture patterns and the evolution of the modern data stack. From Kimball methodology to Data Vault, dimensional modeling to streaming architectures, I enjoy exploring different approaches to solving data challenges at scale.

Staying at the forefront of platform technologies like Databricks and understanding how to leverage them effectively is both a professional necessity and a personal passion.

© 2025 Alex Dimmock. All rights reserved.